



Predicting Flight Delays through Machine Learning Classifiers at Scale Final Presentation

W261 Fall 2022 Section 5 Group 4:
Nathan Chiu, Dominic Lim, Raul Merino, Javier Rondon

01

Executive Summary

02

Exploratory Data Analysis

03

Data Lineage

04

Model Pipeline

05

Feature Selection

06

Hyperparameter Tuning

07

Model Results & Discussion

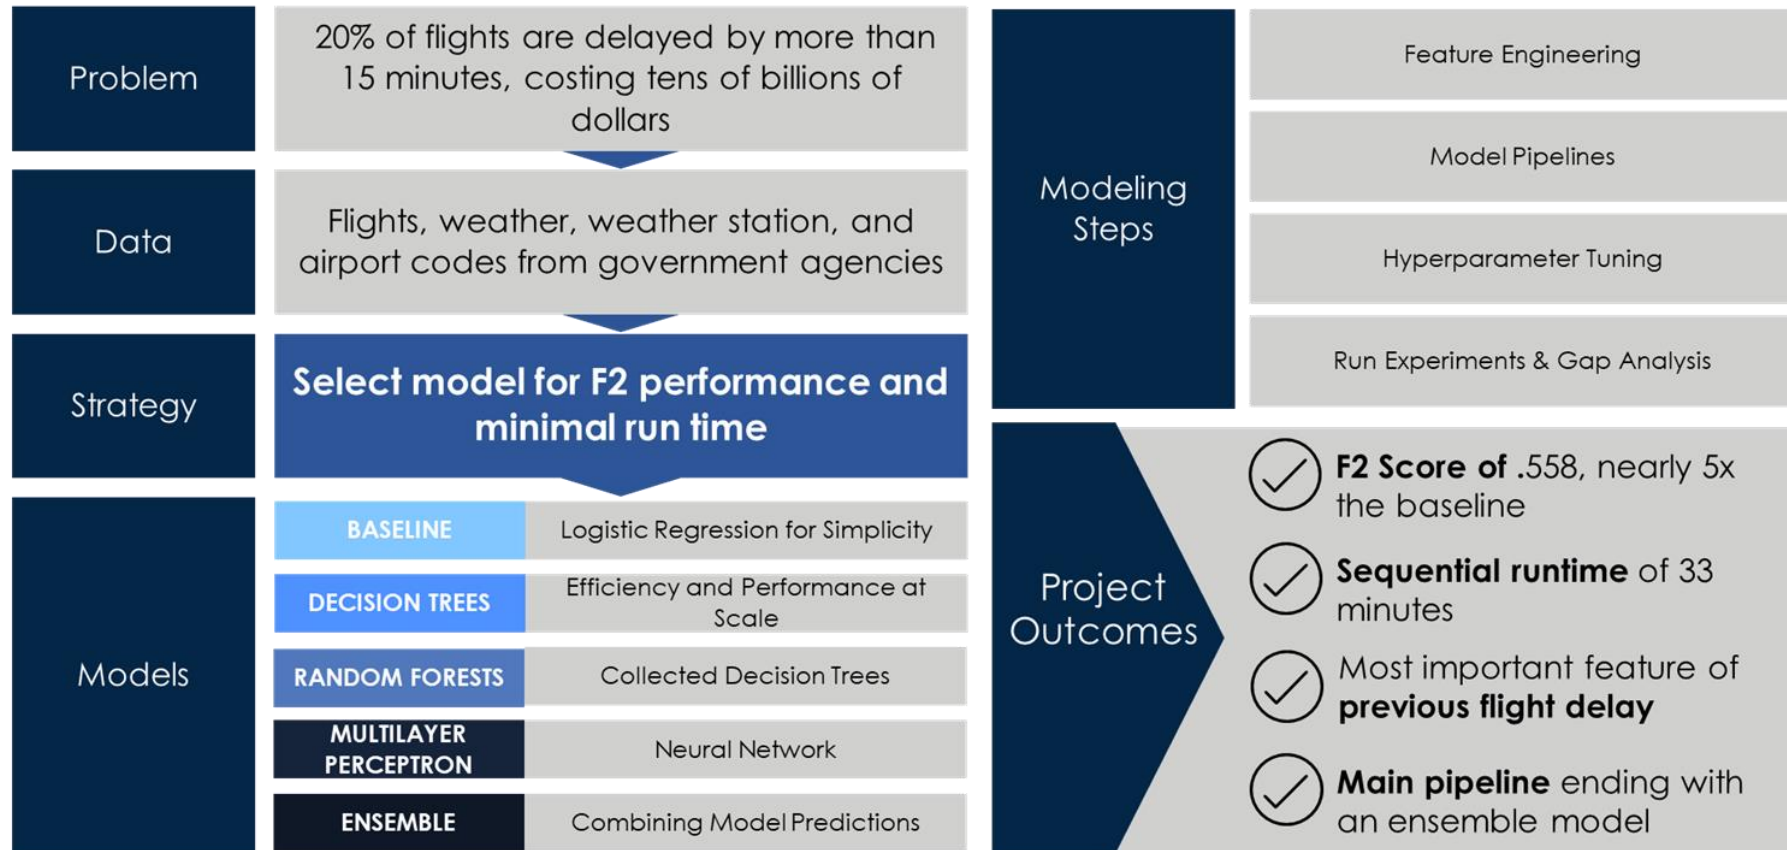
08

Gap Analysis

09

Conclusion

Airlines should implement an **ensemble model** to better predict flight delays for resource allocation/customer service purposes




We used F2 Score as our primary metric of success because we want to focus more on recall to minimize false negatives

What is F2 Score?

$$F_2 = \frac{TP}{TP + 0.2FP + 0.8FN}$$

F2 is the weighted average of precision and recall and we selected the beta value of 2 to focus on recall, which refers to proportion of true positives that are correctly identified

Why use F2 Score?

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP
	Negative (0)	FN 	TN


We are using F2 for more focus on recall to **minimize false negatives**. False negatives occur when the flight is predicted to not be delayed, but ends up being delayed, which is the main driver in the costs in flight delays

With F1, the denominator would contain 0.5 FP and 0.5 FN

We conducted an exploratory data analysis of the flight and weather datasets, focusing on computing % of missing values per feature and understanding the features' distribution, scale, and range of values

Flights Dataset EDA

YEAR	Extreme Weather (%)	NAS - Weather (%)	Late Aircraft - Weather (%)	Total Weather (%)
2015	5.37	10.37	6.87	22.61
2016	4.70	11.26	6.85	22.81
2017	4.59	11.94	7.09	23.62
2018	6.13	12.10	7.92	26.14
2019	5.95	12.52	7.97	26.43
2020	7.86	7.23	4.97	20.06
2021	7.41	6.41	5.30	19.13




- Our EDA suggests that **~20% of total delays** (in minutes) are attributable to weather conditions.

- The increasing average minutes of delay as time goes on suggests **network effects at play** i.e. delays accumulating as the day progresses and planes with delayed flights having their subsequent flights delayed

Data Lineage

We created a data lineage tracing how we turned the raw data files into our joined files and added in features that we plugged into our model pipelines

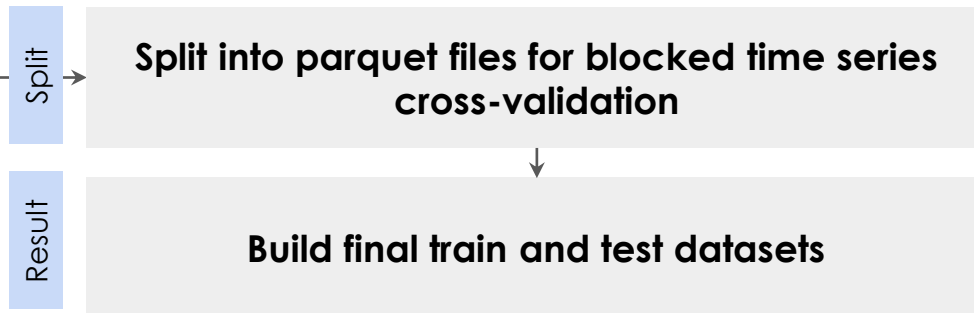
1. Joining Raw Data Files Into Joined Dataset



2. Pulling in Features Created

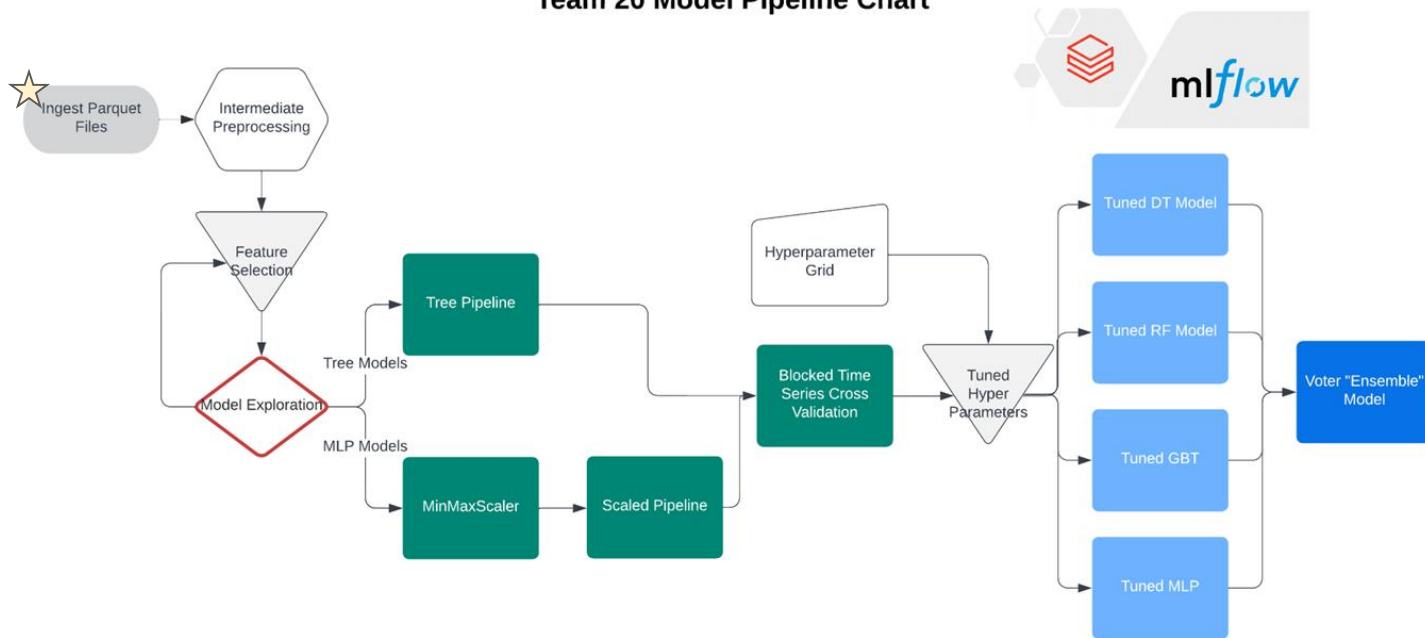


3. Split for Cross-Validation, Train, & Test



We focused on generalizing functions involved in the pipeline to make it easier to adjust parameters and run a multitude of experiments

Team 20 Model Pipeline Chart



* The modeling pipeline begins with ingesting the joined parquet files that were split by FL_DATE for our Blocked Time Series Cross Validation

Feature Name	Description
Previous Flight Delay	Airlines have a finite number of aircrafts, so each aircraft has a route that it follows every day, going from airport to airport often involving back to back scheduled flights. An earlier delay may affect subsequent flights for the same aircraft
Pagerank Features	PageRank describes an airport's importance and influence, which can describe how delays are spread throughout a network of airports.
Delay States	The delay state represents the network's delay patterns at a point in time
Weather Features	The categorical features indicate the presence of weather related to flight delays such as thunderstorms, snow, fog and ice
Average Airport Delay	We created a feature for the percentage of flights that are delayed in a given time window
Airport Capacity	The ratio of actual flights that depart over scheduled flights out of an airport

We conducted an exploratory data analysis of the newly engineered features, focusing on understanding the features' distribution, scale, and range of values

Conceptual Visualization



What and why use Previous Flight Delay?

- **What:** We built the previous delay by tracking the tail number of the plane and **tagging whether or not the plane's previous flight was delayed or not**
 - E.g. Plane ABC is flying from Canada to Eastern Europe, but the previous flight from Australia to Canada was delayed would be tagged as being previously delay i.e. has a value of 1
- **Why:** After an initial analysis, we saw that this had a relatively **high correlation (> 0.3)** with the current's flight delay
 - When conducting our initial EDA we thought that past flights would bear an impact on current flights so we decided to build this feature to test this hypothesis

We conducted an exploratory data analysis of the newly engineered features, focusing on understanding the features' distribution, scale, and range of values

Delay State

Delay patterns in 2015


Cluster 1

Cluster 4

Cluster 5

Cluster 2


Cluster 3



- In the delay state cluster with the most delays stem from flights that involve **DFW, ORD and LAX**

Airport PageRank

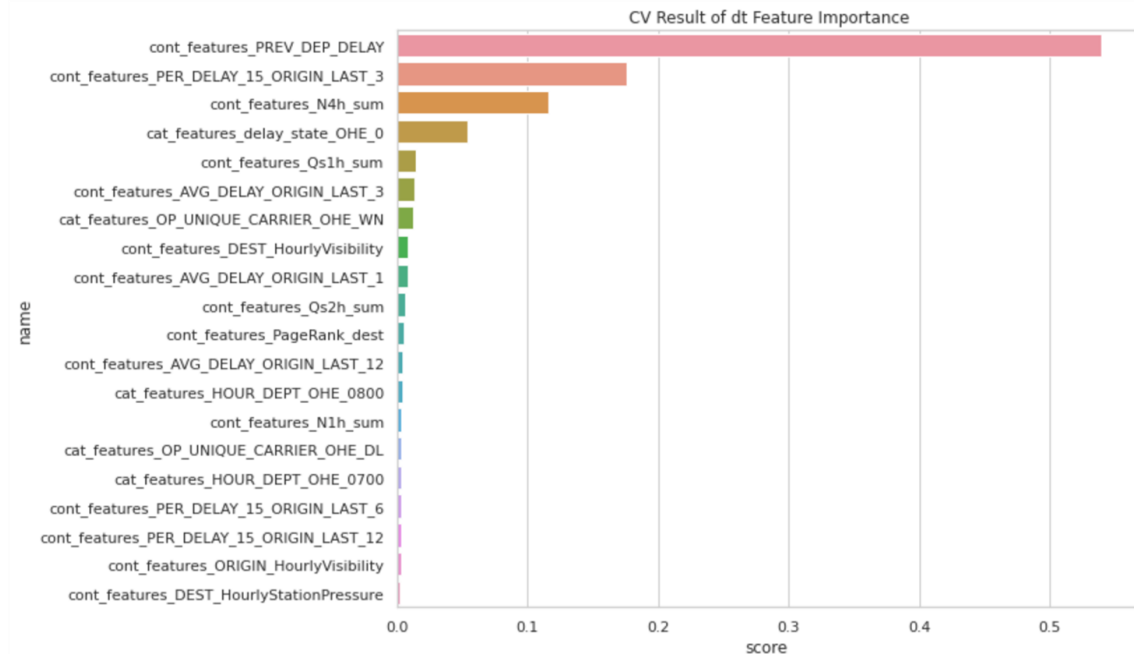
Airport PageRank 2015-2021



- PageRank shows that the most important airports are **ATL and DFW** and changes in rank across time

We ran decision tree models with different categories of features and select top three features per category by feature importance

Determining Feature Importance Against CV Data



We selected the three most important features by feature category. Notable high performing features include:

- Previous Delay (PREV_DEP_DELAY)
- PER_DELAY_15_ORIGIN_LAST_3

Feature categories:

- Weather features
- Pagerank features
- Airport capacity
- Delay state
- Previous flight delay
- Other flight features

We utilized **feature importance**, a measure of the decrease in **node purity** weighted by the **probability of reaching that node** to score and rank features above


Experimental Results

We compared primary metrics of success like F2 across tuned models and the Ensemble models performed the best

Experiment Results with Hyperparameters

Model	Layers	Max Bins	Max Depth	Max Iterations	Number of Trees	Train F2	Train ROC AUC	Train Precision	Train Recall	Test F2	Test ROC AUC	Test Precision	Test Recall
MLP	[44, 44, 2]	-	-	100	-	0.641	0.748	0.716	0.619	0.519	0.755	0.388	0.589
Decision Tree	-	350	10	-	-	0.617	0.765	0.760	0.589	0.540	0.764	0.411	0.586
Gradient Boosted Tree	-	100	10	6	-	0.630	0.772	0.756	0.605	0.546	0.771	0.405	0.599
Random Forest	-	50	10	-	100	0.642	0.765	0.737	0.622	0.547	0.765	0.384	0.612
Ensemble	-	-	-	-	-	-	-	-	-	0.558	-	0.366	0.643

Model Results Bar Chart



Voting Mechanism

***One Positive Voting:** If one model suggests delay, predict DELAY


Vote by Majority: The majority prediction of DELAY or NO DELAY

One Negative Voting: If one model suggests no-delay, predict NO DELAY

* Voting mechanism we ultimately selected for the ensemble model

We also wanted to implement novel approaches including the use of ensemble methods whereby all four models (hyper-parameterized Decision Tree, Random Forest, Gradient Boosted Tree, and Multilayer Perceptron) "vote" on the final prediction


Ensemble (Best Model) Confusion Matrix



Metrics:

- **F2 Score: .558**
- **Precision: .366**
- **Recall is: .643**

Distribution of Flight Predictions by Models in Agreement



How to read chart: For 300K, none of the models correctly predicted them as delayed. For 500K flights, all models correctly predicted them as delayed

We compared primary metrics of success like F2 across tuned models and the Ensemble models performed the best

Experiment Wins

1. **Model Novelty:** Experimented with ensemble model with varying voting mechanisms and found the positive voting drove the **better recall and F2 metric**
2. **Scientific Approach to Feature Selection:** Used relative importances i.e. key decision boundaries of the decision trees

Experiment Areas of Opportunity / Next Steps

1. **Altering Pipeline Inputs:** Will explore ensemble model on other models besides the four specified earlier. Models with different inputs i.e. only weather features
2. **Pull in Additional External Datasets:** We explored several external datasets like airline ratings and natural disasters that we would like to join in if given more time
3. **Further Investigate Poorly Predicted Flights:** Some flights were not predicted well by any of our models, which was also true for our training data. We can isolate these data points and create models for this data

We performed a Gap Analysis comparing our best model results with models from peers and then identified strengths and opportunities for improvement to incorporate in our models

Our Best Model Performance

.558

F2 Score

Comparable Best Model Performance

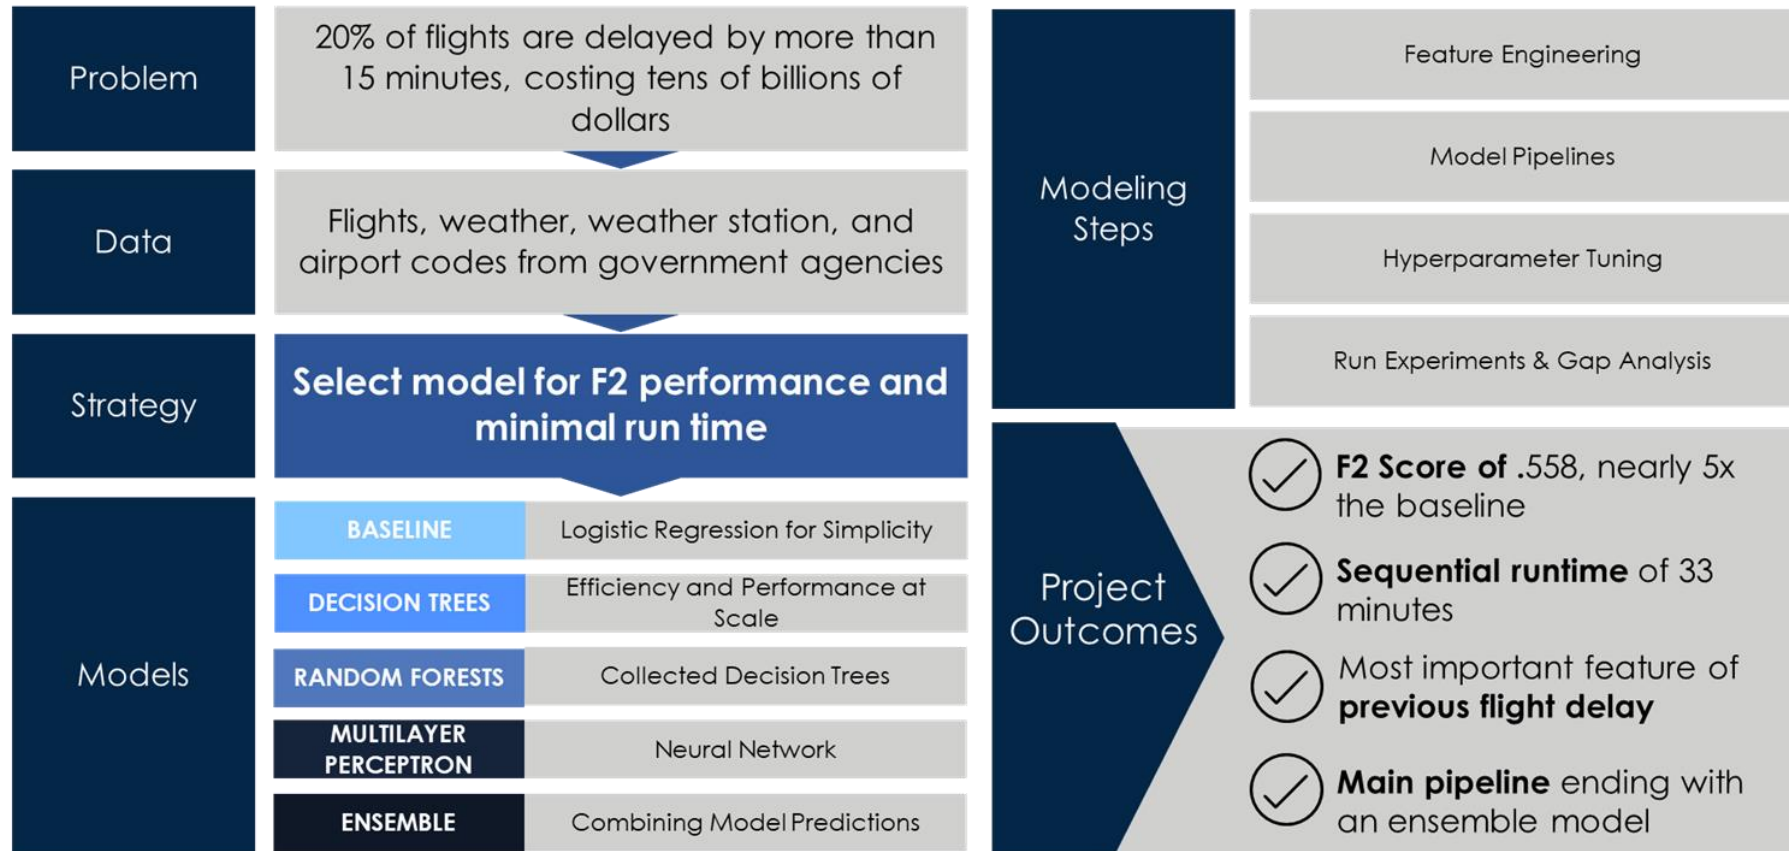
.750

F2 Score

- One particular method that we did not explore was **weighing our training data differently**. We could have determined if a particular year was more relevant than others for the 2021 data, and weighed that year higher than the rest
- Another feature that we did not focus as much as we would have liked was **departure and arrival times**. We could have turned that "continuous" feature into a categorical ones. It's unlikely the relationship with the delay was linear, instead different times of day could have had a different result (as shown in our [initial EDA](#))
- Since we explored a variety of different models, it doesn't seem that the gap came from not trying a particular model. We could have however done **more extensive hyperparameter tuning and more experimentation**

Conclusion

Airlines should implement an **ensemble model** to better predict flight delays for resource allocation/customer service purposes



Main Deck

[Executive Summary](#)

[Project Metrics](#)

[Raw Data EDA](#)

[Data Lineage](#)

[Model Pipeline](#)

[Feature Engineering](#)

[Previous Delay Feature EDA](#)

[Delay State & Pagerank Feature Engineering EDA](#)

[Feature Importance](#)

[Ensemble Model Overview](#)

[Model Results](#)

[Model Discussions \(Wins & Opportunities\)](#)

[Gap Analysis](#)

[Conclusion](#)

Appendix

[Data Lineage Code](#)

[Model Pipeline Code](#)


[Model Pipeline Description](#)

[What is a MLP?](#)

[Weather Dataset EDA](#)

[Decision Tree F2 Scores](#)

[Random Forest F2 Scores](#)

A photograph of an airplane wing, seen from a passenger's perspective, flying over a vast expanse of white clouds. The sky is a mix of blue and orange, suggesting a sunset or sunrise. The wing is white with dark structural lines. The word "Appendix" is written in a black, sans-serif font, centered horizontally and partially overlaid by a semi-transparent white band.

Appendix

We created a data lineage tracing how we turned the raw data files into our joined files and added in features that we plugged into our model pipelines

Composite Key Code

Create composite key to join weather and flights dataset

```
def create_composite_key(code, timestamp):
    return f'{code}_{timestamp}'

create_composite_key = udf(create_composite_key)

# Create composite key for both datasets
df_weather_icao_needed_tz = df_weather_icao_needed_tz.withColumn("CODE_STATION_TIMESTAMP", create_composite_key('icao', 'HOUR_TIMESTAMP'))
flights_icao_tz = flights_icao_tz.withColumn("CODE_TIMESTAMP", create_composite_key('icao', 'HOUR_WEATHER_TIMESTAMP')) \
    .withColumn("TWO_CODE_TIMESTAMP", create_composite_key('icao', 'TWO_HOUR_WEATHER_TIMESTAMP')) \
    .withColumn("THREE_CODE_TIMESTAMP", create_composite_key('icao', 'THREE_HOUR_WEATHER_TIMESTAMP'))

# Update the datasets
flights_icao_tz.write.mode("overwrite").parquet(f"{blob_url}/flights_with_icao_tz")
df_weather_icao_needed_tz.write.mode("overwrite").parquet(f"{blob_url}/weather_with_icao_tz")

# Load the new datasets
flights_icao_tz = spark.read.parquet(f"{blob_url}/flights_with_icao_tz") # 42430592
df_weather_icao_needed_tz = spark.read.parquet(f"{blob_url}/weather_with_icao_tz") # 1057832
print(flights_icao_tz.count())

# Drop unnecessary columns
columns_to_drop = ['airport_name', 'airport_city', 'airport_country', 'airport_tz', 'year', 'airport_subd', 'country', 'elevation', 'iata', 'airport_lon', 'airport_lat', 'icao']
df_weather_icao_needed_tz = df_weather_icao_needed_tz.drop(*columns_to_drop)
```

We focused on generalizing functions involved in the pipeline to make it easier to adjust parameters and run a multitude of experiments

Model Pipeline Reusable Functions

General Functions

```

1 def read_clean(parquet_string):
2     dataset = spark.read.parquet(f"{blob_url}/{parquet_string}")
3
4     # Make sure the target variable is not null
5
6     dataset = dataset.where("DEP_DELAY is not NULL")
7
8     dataset = dataset.withColumn("PREV_DEP_DELAY", col("PREV_DEP_DELAY").cast('int'))
9
10    dataset= dataset.withColumnRenamed("DEP_DELAY", "label")
11    dataset= dataset.withColumn("HOUR_DEPT", substring("DEP_TIME",1,4))
12
13    for col_name in cont_feats:
14        dataset = dataset.withColumn(col_name, col(col_name).cast('float'))
15
16    dataset = dataset.na.drop(subset=["ORIGIN_HourlyStationPressure",
17                                     "DEST_HourlyStationPressure",
18                                     "ORIGIN_HourlyDryBulbTemperature",
19                                     "DEST_HourlyDryBulbTemperature",
20                                     "ORIGIN_HourlyVisiblity",
21                                     "DEST_HourlyVisiblity"])
22    .fillna(0, subset=["ORIGIN_HourlyPrecipitation",
23                      "ORIGIN_HourlyWindDirection",
24                      "ORIGIN_HourlyWindSpeed",
25                      "DEST_HourlyPrecipitation",
26                      "DEST_HourlyWindDirection",
27                      "DEST_HourlyWindSpeed"])
28
29    # PLACEHOLDER TO DEAL WITH NULLS/NAs for QRN Features
30
31    return dataset
32
33
34    def create_parameters(parameter_grid):
35        param_names = list(parameter_grid.keys())
36        param_values = parameter_grid.values()
37        combinations = list(itertools.product(*param_values))
38        return (param_names, combinations)
39
40    def downsample(train_df):
41        n_delays = train_df.filter(f.col("label") == 1).count()
42        n_no_delays = train_df.filter(f.col("label") == 0).count()
43
44        total = n_delays + n_no_delays
45        keep_percent = n_delays / n_no_delays
46
47        train_delay = train_df.filter(f.col("label") == 1)
48        train_non_delay = train_df.filter(f.col("label") == 0).sample(withReplacement=False, fraction=keep_percent, seed=741)
49        train_downsampled = train_delay.union(train_non_delay)
50        return train_downsampled
51
52    def get_metrics(pred_df):
53        preds_mc_rdd = pred_df.select(["prediction", "label"]).rdd
54        preds_b_rdd = pred_df.select(["label", "probability"]).rdd.map(lambda row: (float(row["probability"])[1], float(row["label"])))
55        metrics_mc = MulticlassMetrics(preds_mc_rdd)
56        metrics_b = BinaryClassificationMetrics(preds_b_rdd)
57        F2 = np.round(metrics_mc.fMeasure(label=1.0, beta=2.0), 4)
58        au_ROC = metrics_b.areaUnderROC
59        return F2, au_ROC

```

Model Pipeline Code

Specify Model Pipelines

```

1 def tree_pipeline(model):
2
3     """Pipeline for tree models - DT, RF, GBT"""
4
5     assembler_cont = VectorAssembler(inputCols=cont_feats,
6                                     outputCol="cont_features")
7
8     #Categorical Features that need to be binned before being used as categorical variables
9     # columns_to_bucketize = ['wind_dir_avg_origin', 'wind_dir_avg_dest']
10    # splits = [[1 for i in range(0,361,20)], [1 for i in range(0,361,20)]]
11    # bucketizer = Bucketizer(splitsArray=splits, inputCols=columns_to_bucketize, outputCols=[c+"_bucketized" for c in columns_to_bucketize])
12
13    indexer = StringIndexer(inputCols=columns_categorical,
14                            outputCols=[c+"_indexed" for c in columns_categorical]).setHandleInvalid("keep")
15
16    ohe = OneHotEncoder(inputCols=[c+"_indexed" for c in columns_categorical],
17                        outputCols=[c+"_OHE" for c in columns_categorical]).setHandleInvalid("keep")
18
19    assembler_categ = VectorAssembler(inputCols=[c+"_OHE" for c in columns_categorical],
20                                     outputCol="cat_features")
21
22    assembler = VectorAssembler(inputCols=["cat_features", "cont_features"],
23                               outputCol="features")
24
25
26    pipeline = Pipeline(stages=[assembler_cont, #bucketizer,
27                                indexer, ohe, assembler_categ, assembler, model])
28
29    return pipeline

```

Train Model (Decision Tree)

```

dt_best_parameters, dt_best_score = blockTimeSeriesCV(parquet_string='full',
                                                       sampling = 'down',
                                                       param_grid = paramGrid_dt,
                                                       pipeline_fn = scaled_pipeline,
                                                       model_type='dt',
                                                       k=2,
                                                       metric='f2'
                                                       )

dtModel, pred_result_dt = validation(full_train_df, full_test_df,
                                     sampling = 'down',
                                     model_type = 'dt',
                                     best_parameters = dt_best_parameters,
                                     pipeline_fn = scaled_pipeline
                                     )

```

Evaluate Model (RF)

```

1 rfModel, pred_result_rf = validation(full_train_df, full_test_df, sampling = 'down',
2                                   model_type = 'rf',
3                                   best_parameters = rf_best_parameters,
4                                   pipeline_fn = tree_pipeline
5                                   )

```

37 Spark Jobs

CV Result of rf

	Train	Test
ROC AUC	0.715	0.712
F2 Score	0.537	0.469
Recall	0.592	0.598
Precision	0.743	0.375
Accuracy	0.664	0.770

Support for multi-dimensional indexing (e.g., 'obj[:, None]') is deprecated and will be re

We focused on generalizing functions involved in the pipeline to make it easier to adjust parameters and run a multitude of experiments

Feature Selection

We began by running decision tree models with different categories of features:

- Weather Features
- Airport Capacity (QRN)
- Airport PageRank
- Clustered Delay States
- Previous Flight Feature (based on Tail Number)
- Other Flight Features (Airline Carrier, Seasonality)

Hyperparameter Tuning

Once features were selected, we experimented with combinations of parameters against cross validation data

- Decision Trees / MLP: VectorAssembler, MinMaxScaler
- Decision Tree Loss Function: Gini Impurity


Model Selection

Once we selected the best hyperparameters, we compared the primary metrics like F2 score, precision, and recall across all models:

- Used average F2 score to fit the full train dataset and evaluate the full test dataset

We implemented a multilayer perceptron (MLP), a fully connected neural network easily integrated into our model pipeline


What is a Multilayer Perceptron and why use it?



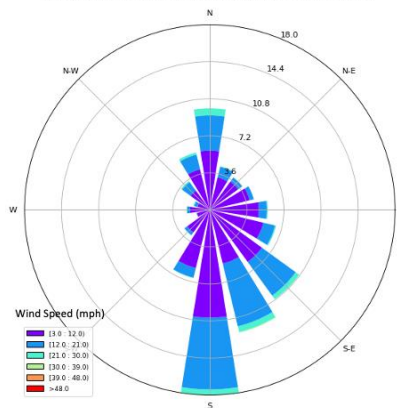
- The MLP utilizes an initial layer of m nodes, representing the number of features, with a final output layer of 2
- We had limited time to experiment with hyperparameter tuning, where we primarily changed the number of layers and the number of nodes per layer
- We ultimately found that a MLP architecture of (44 - Sigmoid - 44 - Sigmoid - 2 - Softmax) produced the best F2 score for our selected features.

We conducted an exploratory data analysis of the flight and weather datasets, focusing on computing % of missing values per feature and understanding the features' distribution, scale, and range of values


Weather Dataset EDA



Distribution of Hourly Wind Direction in DFW 2015



DFW Airport Runway Diagram




For informational use. Effective March 05, 2015 – April 02, 2015.


- The distributions of the temperature, snowfall, rain, and percentage of delayed flights suggest that weather feature are highly correlated to certain flight delays

- Features such as wind speed and direction were reviewed to assess their usefulness to explain flight delays.

F2 Score by Decision Tree Model



Random Forest Overall Performance



A photograph taken from a high altitude, looking out from an airplane window. The wing of the aircraft is visible in the foreground, extending from the bottom left towards the top right. The wing is white with dark structural lines. Below the wing, a vast expanse of white, fluffy clouds stretches across the horizon. The sky above the clouds is a mix of blue and orange, suggesting a sunset or sunrise. The word "Graveyard" is written in a large, black, sans-serif font, centered horizontally across the middle of the image, partially overlapping the wing and the sky.

Graveyard